

# Efficient Calculation of the Genomic Relationship Matrix

Martin Schlather\*

Correspondence:

schlather@math.uni-mannheim.de  
Universität Mannheim, Institut für  
Mathematik, C8, 5, 68131  
Mannheim, Germany

Full list of author information is  
available at the end of the article

\*Equal contributor

## Abstract

**Background:** The calculation of a genomic relationship matrix needs a large number of arithmetic operations. Therefore, fast implementations are of interest. The currently fastest implementations use AVX floating-point arithmetics.

**Results:** Our fastest algorithm is more accurate and  $25\times$  faster than a AVX double precision floating-point implementation.

**Conclusions:** The spectrum of presented methods suggests that further improvement might be possible and that bit manipulation in combination with hash tables might be of relevance also for other calculation problems.

**Keywords:** crossproduct; covariance matrix; genetic relationship; genomic relatedness matrix; hash table

## Background

The genomic relationship matrix (GRM) is the covariance matrix calculated from the SNP information of the individuals, i.e., from the minor allele counts [1]. It is an important ingredient in mixed models and generalized mixed models for analyses and predictions in genetics [2].

Let  $n$  be the number of individuals and  $s$  the number of SNPs per individual. Then, the calculation of the GRM needs of order  $sn^2$  arithmetic operations. Most software packages use floating-point arithmetics, for instance the R packages `AGHmatrix` [3], `qgg` [4], `rrBLUP` [5], `snpReady` [6], and `GENESIS` [7]. The software `GCTA` [8] treats missings explicitly. The package `SNPRelate` [9] also uses floating-point arithmetics for the covariance matrix, but uses bit manipulation algorithms for other calcula-

tions (identity-by-descent estimates). PLINK [10] profits from bit manipulations for<sup>1</sup>  
calculating the uncentred covariance matrix.<sup>2</sup>

<sup>3</sup> In this paper we present a couple of ideas, how the GRM can be calculated<sup>3</sup>  
<sup>4</sup>efficiently from a SNP matrix. We assume that no values are missing. The emphasis<sup>4</sup>  
<sup>5</sup>will be on algorithms that allow for a vectorized implementation (SIMD) and that<sup>5</sup>  
<sup>6</sup>take into account that the entries of the SNP matrix are most efficiently coded by<sup>6</sup>  
<sup>7</sup>2 bits, namely for the values 0, 1 and 2, as they are present in diploid organism<sup>7</sup>  
<sup>8</sup>under the common assumption of biallelic markers.<sup>8</sup>

<sup>9</sup>

## <sup>10</sup>Results<sup>10</sup>

<sup>11</sup>The standard mathematical formula for the GRM requires floating point arith-<sup>11</sup>  
<sup>12</sup>metics. An algebraic reformulation shows that the cost intensive part involves only<sup>12</sup>  
<sup>13</sup>integers. Since the most elementary numbers need only a minimum of 2 bits, a<sup>13</sup>  
<sup>14</sup>diversity of approaches for the integer arithmetics is thinkable. The investigated<sup>14</sup>  
<sup>15</sup>methods are in brief:<sup>15</sup>

- <sup>16</sup> • **Multiply** : uses 16-bit arithmetics<sup>16</sup>
- <sup>17</sup> • **Hamming2** : uses pop counts (the number of bits that are 1)<sup>17</sup>
- <sup>18</sup> • **ThreeBit** : uses a 3-bit representation and a single large hash table<sup>18</sup>
- <sup>19</sup> • **TwoBit** : uses the 2-bit representation and two large hash tables<sup>19</sup>
- <sup>20</sup> • **Shuffle** : similar to **TwoBit**, but with two tiny hash tables<sup>20</sup>
- <sup>21</sup> • **Packed** : uses 4-bit arithmetics<sup>21</sup>

<sup>22</sup>They are all available through **crossprodx** in the package **miraculix** [11]. Tables 1<sup>22</sup>  
<sup>23</sup>and 2 show that **Shuffle** is the fastest method, which is, in the SSSE3 implemen-<sup>23</sup>  
<sup>24</sup>tation, 35× faster than **crossprod** of R [12].<sup>24</sup>

<sup>25</sup>**Table 1** Acceleration of the calculation of  $M^T M$  by **crossprodx**.<sup>25</sup>

	Shuffle	Packed	Hamming2	Multiply	ThreeBit	TwoBit
acceleration	35×	28×	24×	17×	17×	15×
SIMD	SSSE3	SSE2	SSE2	SSE2	SSE2	none

The reference point is **crossprod** in R.

<sup>29</sup>**Table 2** Acceleration of the calculation of  $M^T M$  by **AVX2** implementations in **crossprodx**.<sup>29</sup>

	Shuffle	Packed	Multiply	AVX (double)	AVX2 (32-bit integer)
acceleration	48 ×	36×	29×	1.8×	4×

The reference point is **crossprod** in R.

<sup>32</sup> The command **relationshipMatrix** in **miraculix** [11] for calculating the GRM is<sup>32</sup>  
<sup>33</sup> only negligibly slower than **crossprodx**. The AVX2 variant of **Shuffle** is even 48×<sup>33</sup>

<sup>1</sup>faster than `crossprod` [12] and  $48/1.8 \approx 25\times$  faster than a standard AVX double<sup>1</sup>

<sup>2</sup>precision implementation for calculating the crossproduct of an arbitrary matrix,<sup>2</sup>

<sup>3</sup>cf. Table 2. Furthermore, our algorithms have not even any cumulative rounding<sup>3</sup>

<sup>4</sup>error. 4

<sup>5</sup> Tables 1 and 2 also show that the AVX2 performance is hard to predict from the<sup>5</sup>

<sup>6</sup>SSE performance. AVX2 variants for `TwoBit` and `ThreeBit` are not given since full<sup>6</sup>

<sup>7</sup>vectorization is not possible. `Hamming2` has not been pursued because of its memory<sup>7</sup>

<sup>8</sup>demand. 8

<sup>9</sup> For the benchmarks, we used an  $s \times n$  SNP matrix with  $n = 1000$  individuals<sup>9</sup>

<sup>10</sup>and  $s = 5 \cdot 10^5$  SNPs. The calculations were performed on an Intel(R) Core(TM)<sup>10</sup>

<sup>11</sup>i7-8550U CPU @ 1.80GHz with R version 3.6.0 on Xubuntu. Although the code in<sup>11</sup>

<sup>12</sup>`miraculix` is parallelized, we used only a single core for the benchmarks. Nonethe-<sup>12</sup>

<sup>13</sup>less, the AVX2 variant of `Shuffle` takes not more than 7 seconds. 13

<sup>14</sup> The code for the benchmarks is available from the man page of `crossprod` in<sup>14</sup>

<sup>15</sup>`miraculix`. 15

<sup>16</sup> 16

## <sup>17</sup>**Discussion** 17

<sup>18</sup>First, with respect to the memory needs of the SNP matrix, algorithms that use<sup>18</sup>

<sup>19</sup>the 2-bit representation of a SNP value should be preferred. Among them, we have<sup>19</sup>

<sup>20</sup>a sequence of distinct algorithms that differ in their speed-up and their SIMD<sup>20</sup>

<sup>21</sup>requirements: `TwoBit` (15 $\times$ ; SIMD not used), `Packed` (28 $\times$ ; SSE2); `Shuffle` (35 $\times$ ;<sup>21</sup>

<sup>22</sup>SSSE3). 22

<sup>23</sup> Second, the use of perfect hash tables to cut calculations short might be of general<sup>23</sup>

<sup>24</sup>importance. 24

<sup>25</sup> Third, since loading from non-aligned memory allocation is reported to be slower<sup>25</sup>

<sup>26</sup>[13], the package `miraculix` was designed to avoid non-aligned loadings. Tests on<sup>26</sup>

<sup>27</sup>the implemented package however show that the running time by non-aligned load-<sup>27</sup>

<sup>28</sup>ings is not reduced for SSE implementations. The speed is reduced by 5 to 10 %<sup>28</sup>

<sup>29</sup>in AVX2 implementations. As the compressed SNP matrix is made available to<sup>29</sup>

<sup>30</sup>the user as an R object and as the memory allocation by R is only 32-bit aligned,<sup>30</sup>

<sup>31</sup>additional memory is allocated and the SNP matrix is aligned to 128 or 256 bits.<sup>31</sup>

<sup>32</sup> Furthermore, additional zeros are appended so that the virtual number of SNPs is a<sup>32</sup>

<sup>33</sup>multiple of the number of the SNPs that can be treated in a single step. The storing<sup>33</sup>

<sup>1</sup>formats of `TwoBit`, `Multiply`, `Packed`, and `Shuffle`, including their AVX2 variants,<sup>1</sup>  
<sup>2</sup>were made compatible, i.e., the allocations are all based on a 256-bit alignment. A<sup>2</sup>  
<sup>3</sup>check and a reallocation are implemented for the case that memory is moved. This<sup>3</sup>  
<sup>4</sup>might happen when the garbage collector `gc` is called by `R`, for instance. <sup>4</sup>

## <sup>6</sup>Conclusion <sup>6</sup>

<sup>7</sup>The combination of algebraic reformulation, bit manipulations and hash tables can<sup>7</sup>  
<sup>8</sup>reduce largely the computing time on SNP data. In the case of calculating the GRM,<sup>8</sup>  
<sup>9</sup>the computing time could be reduced by factor 25 in comparison to a straightfor-<sup>9</sup>  
<sup>10</sup>ward AVX double precision implementation. As a spectrum of implementations<sup>10</sup>  
<sup>11</sup>exist, there is a chance of further improvement and of further applications of the<sup>11</sup>  
<sup>12</sup>underlying ideas. <sup>12</sup>

## <sup>13</sup>Methods <sup>13</sup>

<sup>14</sup>Let  $M$  be an  $s \times n$  SNP matrix of  $n$  individuals and  $s$  SNPs. We need to consider only<sup>15</sup>  
<sup>16</sup>the fast calculation of the crossproduct  $M^\top M$ , since the GRM  $A$  can be calculated<sup>16</sup>  
<sup>17</sup>from  $M^\top M$  at low costs. This can be seen as follows. <sup>17</sup>

<sup>18</sup>Let  $\mathbf{1}_k$  be the vector of length  $k$  whose components are all equal to 1. The centred<sup>18</sup>  
<sup>19</sup>and normalized GRM  $A$  is calculated as <sup>19</sup>

$$A = (M - P)^\top (M - P) / \sigma^2$$

<sup>22</sup>where <sup>22</sup>

$$P = p \mathbf{1}_n^\top \quad \text{with} \quad p = \frac{1}{n} M \mathbf{1}_n.$$

<sup>25</sup>and <sup>25</sup>

$$\sigma^2 = \sum_{i=1}^s p_i (1 - p_i / 2) \quad \text{with} \quad p = (p_1, \dots, p_s).$$

<sup>28</sup>Note that replacing the value  $p_i$  by the allele frequency  $\tilde{p}_i = p_i / 2$ , we have the usual<sup>28</sup>  
<sup>29</sup>formula for  $\sigma^2$ , <sup>29</sup>

$$\sigma^2 = 2 \sum_{i=1}^s \tilde{p}_i (1 - \tilde{p}_i).$$

<sup>32</sup>Let  $B = M^\top M \mathbf{1}_n$ . Then <sup>32</sup>

$$n^2 \sigma^2 A = n^2 M^\top M - n \mathbf{1}_n B^\top - n B \mathbf{1}_n^\top + \mathbf{1}_n \mathbf{1}_n^\top B \mathbf{1}_n^\top.$$

<sup>1</sup>Hence, the integer-valued matrix  $n^2\sigma^2A$  can easily be calculated from the matrix<sup>1</sup>

<sup>2</sup> $M^\top M$  without any numerical error and at low computational costs of order  $n^2$ .<sup>2</sup>

<sup>3</sup>Now,<sup>3</sup>

$$2n^2\sigma^2 = 2n^2 \sum_{i=1}^s p_i - n^2 \sum_{i=1}^s p_i^2 = 2n\mathbf{1}_s^\top M\mathbf{1}_n - \mathbf{1}_n^\top B.$$

<sup>6</sup>Again,  $2n^2\sigma^2$  can easily be calculated from  $M$  and  $M^\top M$  without any numerical<sup>6</sup>

<sup>7</sup>error. The computational costs are of order  $n(n+s)$ , hence still some magnitudes<sup>7</sup>

<sup>8</sup>smaller than the costs of calculating the crossproduct  $M^\top M$ .<sup>8</sup>

#### <sup>10</sup>Algorithms for scalar products<sup>10</sup>

<sup>11</sup>Instead of considering the crossproduct  $M^\top M$ , it suffices to consider the scalar<sup>11</sup>

<sup>12</sup>product of two vectors  $a = (a_1, \dots, a_s)$  and  $b = (b_1, \dots, b_s)$  whose components  $a_i$ <sup>12</sup>

<sup>13</sup>and  $b_i$  have the values 0, 1 or 2. For simplicity and clarity, we will primarily refer<sup>13</sup>

<sup>14</sup>to SSE commands in the following, and not to AVX.<sup>14</sup>

<sup>15</sup><sup>15</sup>

#### <sup>16</sup>*Simple Multiplication*<sup>16</sup>

<sup>17</sup>An immediate way of calculating the scalar product from a compressed 2-bit rep-<sup>17</sup>

<sup>18</sup>resentation is to extract the first two bits of each of the two vectors  $a$  and  $b$  and<sup>18</sup>

<sup>19</sup>to continue with integer arithmetics. Then the next two bits are extracted using<sup>19</sup>

<sup>20</sup>shifting, and so on. Clearly, this procedure can be vectorized. Of particular advan-<sup>20</sup>

<sup>21</sup>tage here is the SSE2 command `_mm_madd_epi16`, which multiplies and adds two<sup>21</sup>

<sup>22</sup>consecutive 16-bit integers so that only 7 shifts are necessary for a vector of 64 SNP<sup>22</sup>

<sup>23</sup>values, i.e., for 128 bits. We call this method `Multiply`.<sup>23</sup>

#### <sup>24</sup>*Hamming Distance*<sup>24</sup>

<sup>25</sup>The algorithm used in PLINK [14, 10] is based on the idea that a value is represented<sup>25</sup>

<sup>26</sup>by the number of bits that equal 1 in a 4-bit representation. The values of the vectors<sup>26</sup>

<sup>27</sup> $a$  and  $b$  must be coded asymmetrically by two mappings  $f_a$  and  $f_b$ , say, as a coding<sup>27</sup>

<sup>28</sup>by a single mapping is not possible. Then, the bitwise  $\&$ -operator is applied before<sup>28</sup>

<sup>29</sup>the number of 1's is counted. Table 3 gives a possible realisation.<sup>29</sup>

<sup>30</sup><sup>30</sup>

**Table 3 Table of values for the Hamming distance method.**

<sup>31</sup> $f_a(\cdot) \& f_b(\cdot)$	<sup>31</sup> $f_b(0) = 0000_b$	<sup>31</sup> $f_b(1) = 0011_b$	<sup>31</sup> $f_b(2) = 1111_b$	<sup>31</sup>
<sup>32</sup> $f_a(0) = 0000_b$	<sup>32</sup> 0000 <sub>b</sub>	<sup>32</sup> 0000 <sub>b</sub>	<sup>32</sup> 0000 <sub>b</sub>	<sup>32</sup>
<sup>33</sup> $f_a(1) = 0110_b$	<sup>33</sup> 0000 <sub>b</sub>	<sup>33</sup> 0010 <sub>b</sub>	<sup>33</sup> 0110 <sub>b</sub>	<sup>33</sup>
<sup>33</sup> $f_a(2) = 1111_b$	<sup>33</sup> 0000 <sub>b</sub>	<sup>33</sup> 0011 <sub>b</sub>	<sup>33</sup> 1111 <sub>b</sub>	<sup>33</sup>

<sup>1</sup> The number of bits that equal 1 can be calculated by SSE2 commands based<sup>1</sup>  
<sup>2</sup>on work by [15, 16]. We call the method **Hamming2**. (An SSSE3 implementation in<sup>2</sup>  
<sup>3</sup>**miraculix** [11] is called **Hamming3**.) The method can be turned into a particularly<sup>3</sup>  
<sup>4</sup>fast implementation when novel AVX512 commands are used for the pop counts,<sup>4</sup>  
<sup>5</sup>e.g. `_mm512_popcnt_epi64`. See also **SNPRelate** [9] for pop count implementations.<sup>5</sup>  
<sup>6</sup>Still, the storing costs of the SNP matrix  $M$  remain high, namely  $2 \times 4 = 8$  bits<sup>6</sup>  
<sup>7</sup>per SNP in a standard implementation. 7

<sup>8</sup> 8  
<sup>9</sup>*Perfect Hash Table* 9

<sup>10</sup>Let us consider the product of the first elements  $a_1$  and  $b_1$  of the two SNP vec-<sup>10</sup>  
<sup>11</sup>tors  $a$  and  $b$ . Let us code the SNP values 0, 1 and 2 by 3 bits, e.g. as  $000_b$ ,  $011_b$ <sup>11</sup>  
<sup>12</sup>and  $110_b$ , respectively, and denote this mapping by  $f$ . Then, a perfect hash table<sup>12</sup>  
<sup>13</sup>for  $f(a_1)$  &  $f(b_1)$  returns 0, 1, 2, 4 for  $000_b$ ,  $011_b$ ,  $010_b$ , and  $110_b$ , respectively,<sup>13</sup>  
<sup>14</sup>cf. Table 4. 14

<sup>15</sup>**Table 4 Table of values for the ThreeBit method.** 15

<sup>16</sup> $f(\cdot) \& f(\cdot)$	$f(0) = 000_b$	$f(1) = 011_b$	$f(2) = 110_b$	<sup>16</sup>
$f(0) = 000_b$	$000_b$	$000_b$	$000_b$	<sup>17</sup>
$f(1) = 011_b$	$000_b$	$011_b$	$010_b$	<sup>17</sup>
$f(2) = 110_b$	$000_b$	$010_b$	$110_b$	<sup>18</sup>

<sup>19</sup> Subvectors of length  $k$  of  $a$  and  $b$  can be treated in the same way and the perfect <sup>19</sup>  
<sup>20</sup>hash table returns then the scalar product of the subvectors. The hash table will <sup>20</sup>  
<sup>21</sup>be indexed by  $3k$ -bit numbers, i.e. by values between 0 and  $2^{3k} - 1$ . Since  $k$  should <sup>21</sup>  
<sup>22</sup>be as large as possible at a smallish size of the hash table, and  $3k$  bits should fit <sup>22</sup>  
<sup>23</sup>nicely into 1, 2 or 4 bytes, the only reasonable choice for  $k$  is  $k = 5$ , so that 15 bits <sup>23</sup>  
<sup>24</sup>in a 16-bit representation of a vector with  $k = 5$  components are used. The precise <sup>24</sup>  
<sup>25</sup>size of the hash table is then  $110\,110\,110\,110\,110_b + 1 = 28087$  bytes. We call this <sup>25</sup>  
<sup>26</sup>method **ThreeBit**. 26

## <sup>28</sup>*Two Perfect Hash Tables* 28

<sup>29</sup>Since we did not find a simple way to use a single hash table based on a 2-bit <sup>29</sup>  
<sup>30</sup>representation of the SNP values, we consider here two hash tables and the two <sup>30</sup>  
<sup>31</sup>bitwise operators  $\&$  and  $|$ . The first hash table should return 1 and 4 for  $01_b$  and <sup>31</sup>  
<sup>32</sup> $10_b$ , respectively, while the second hash table returns 2 for  $11_b$ . All other values in <sup>32</sup>  
<sup>33</sup>the hash tables are 0, cf. Table 5. 33

**Table 5 Tables of values for the TwoBit method.**

$\&$	$00_b$	$01_b$	$10_b$	$ $	$00_b$	$01_b$	$10_b$
$00_b$	$00_b$	$00_b$	$00_b$	$00_b$	$00_b$	$01_b$	$10_b$
$01_b$	$00_b$	$01_b$	$00_b$	$01_b$	$01_b$	$01_b$	$11_b$
$10_b$	$00_b$	$00_b$	$10_b$	$10_b$	$10_b$	$11_b$	$10_b$

Then, the sum of the two table values yields the product  $a_1b_1$ . Scalar products of subvectors of length  $k$  can also be treated by two hash tables. Since the size of both hash tables is of order  $2^{2k}$ , one possible choice is  $k = 8$ , so that the size of the second hash table is 65536 bytes. We call this method **TwoBit**.

The disadvantage of **TwoBit** (and **ThreeBit**) is that the look-up in the hash table prohibits a full vectorization. A much better choice is therefore  $k = 2$ : the SSE3 command `_mm_shuffle_epi8` looks 16 values up at once in a hash table of size 16. We call this variant **Shuffle**.

#### *Packed arithmetics*

A last idea is to emulate a multiplication by bitwise operations and partial sums.

Let  $\gg$  denote the bitwise shift operator and let

$$c_i = a_i \& b_i,$$

$$d_i = (c_i \gg 1) \& 01_b,$$

$$e_i = (a_i \gg 1) \& b_i \& 01_b, \text{ and}$$

$$f_i = a_i \& (b_i \gg 1) \& 01_b,$$

so that  $c_i + 2d_i = a_ib_i$  if  $a_i = b_i$  and 0 else. Furthermore  $2(e_i + f_i) = a_ib_i$  if  $a_ib_i = 2^{23}$

and 0 else. In total, we have  $c_i + 2d_i + 2(e_i + f_i) = a_ib_i$ . Let  $g_i = d_i + e_i + f_i$ .

Since  $g_i = d_i | e_i | f_i$  for the bitwise operator  $|$ , only the values of  $g_i$  and  $c_i$  need

to be summed up. An immediate extraction of the values of  $g = (g_1, \dots, g_s)$  and

$c = (c_1, \dots, c_s)$  by shifting as in the **Multiply** algorithm would be rather expensive.

Instead, a 4-bit arithmetic can be introduced in an intermediate step for the four vectors

$$\dots 00_b c_3 00_b c_1,$$

$$\dots 00_b c_4 00_b c_2,$$

$$\dots 00_b g_3 00_b g_1, \text{ and}$$

1  $\dots 00_b g_4 00_b g_2,$  1

2 2

3 3

4 which can be obtained by two shifts and 4 bitwise &-operations in total, if the<sup>4</sup>

5 ordering in the memory is  $\dots c_4 c_3 c_2 c_1$ . Since the value of  $c_i$  is at most 2, a sevenfold<sup>5</sup>

6 summation of the first two displayed vectors leaves each component within its 4<sup>6</sup>

7 bits (using any unsigned integer SIMD addition). Since the value of  $g_i$  is either 0<sup>7</sup>

8 or 1, even a fifteen-fold summation is possible for the last two displayed vectors.<sup>8</sup>

9 Afterwards, the 4-bit values are extracted and further summed up. We call this<sup>9</sup>

10 method **Packed**. 10

11 The novel AVX512 command `_mm512_popcnt_epi64` might improve this ap-<sup>11</sup>

12 proach, as it allows to count the number of bits being 1 in  $c_i$ ,  $d_i$  and  $e_i \mid f_i$ , so<sup>12</sup>

13 that the number of products is counted that equal (i)  $1^2$  or  $2^2$ , (ii)  $2^2$ , and (iii)  $1 \cdot 2$ <sup>13</sup>

14 or  $1 \cdot 2$ . This variant can be seen as a 2-bit analogue of the **Hamming2** algorithm,<sup>14</sup>

15 and will be implemented in future. 15

16 16

## 17 **List of abbreviations** 17

18 **AVX**: Advanced vector extensions 18

19 **SIMD**: Single instruction, multiple data 19

20 **SNP**: Single nucleotide polymorphism 20

21 **SSE**: Streaming SIMD extensions 21

22 **SSSE**: Supplemental streaming SIMD extensions 22

23 **GRM**: Genomic relationship matrix 23

24 24

## 25 **Declarations** 25

26 **Ethics approval** 26

27 Not applicable. 27

28 **Consent for publication** 28

29 Not applicable. 29

30 **Availability of the data and materials** 30

31 Not applicable. 31

32 **Competing interests** 32

33 The author declares that he has no competing interests. 33

34 **Funding** 34

35 None. 35



<sup>1</sup>	<b>Author's contributions</b>	1
<sup>2</sup>	Not applicable.	2
<sup>3</sup>	<b>Acknowledgements</b>	3
	The author is grateful to Torsten Pook, Christopher Dörr and Alexander Freudenberg for hints and for comments.	
<sup>4</sup>	<b>References</b>	4
<sup>5</sup>	1. VanRaden, P.M.: Efficient methods to compute genomic predictions. <i>Journal of Dairy Science</i> <b>91</b> (11),	5
<sup>6</sup>	4414–4423 (2008)	6
<sup>7</sup>	2. Walsh, B., Lynch, M.: <i>Evolution and Selection of Quantitative Traits</i> . Oxford University Press, Oxford (2018)	7
<sup>8</sup>	3. Rampazo Amadeu, R., Cellon, C., Olmstead, J.W., Franco Garcia, A., Resende Jr, M.F.R.: AGHmatrix: R	8
	package to construct relationship matrices for autotetraploid and diploid species. <i>The Plant Genome</i> <b>9</b> (3), 1–10	
<sup>9</sup>	(2016)	9
<sup>10</sup>	4. Soerensen, P., Rohde, P.D., Soerensen, I.F.: qgg: An R package for Quantitative Genetic and Genomic	10
<sup>11</sup>	analyses. <a href="http://psoerensen.github.io/qgg/">http://psoerensen.github.io/qgg/</a> (2019)	11
<sup>12</sup>	5. Endelman, J.B.: Ridge regression and other kernels for genomic selection with R package rrBLUP. <i>Plant</i>	12
<sup>13</sup>	<i>Genome</i> <b>4</b> , 250–255 (2011)	13
<sup>14</sup>	6. Granato, I., Fritsche-Neto, R.: snpReady: Preparing Genotypic Datasets in Order to Run Genomic Analysis. R	14
<sup>15</sup>	package version 0.9.6. <a href="https://cran.r-project.org/web/packages/snpReady/index.html">https://cran.r-project.org/web/packages/snpReady/index.html</a> (2018)	15
<sup>16</sup>	7. Gogarten, S.M., Sofer, T., Chen, H., Yu, C., Brody, J.A., Thornton, T.A., Rice, K.M., Conomos, M.P.: Genetic	16
<sup>17</sup>	association testing using the GENESIS R/Bioconductor package. <i>Bioinformatics</i> <b>35</b> , 5346–5348 (2019)	17
<sup>18</sup>	8. Yang, J., Lee, S.H., Goddard, M.E., Visscher, P.M.: GCTA: a tool for genome-wide complex trait analysis. <i>The</i>	18
<sup>19</sup>	<i>American Journal of Human Genetics</i> <b>88</b> (1), 76–82 (2011)	19
<sup>20</sup>	9. Manichaikul, A., Mychaleckyj, J.C., Rich, S.S., Daly, K., Sale, M., Chen, W.-M.: Robust relationship inference	20
<sup>21</sup>	in genome-wide association studies. <i>Bioinformatics</i> <b>26</b> , 2867–2873 (2010)	21
<sup>22</sup>	10. Chang, C.C., Chow, C.C., Tellier, L.C.A.M., Vattikuti, S., Purcell, S.M., Lee, J.J.: Second-generation PLINK:	22
<sup>23</sup>	rising to the challenge of larger and richer datasets. <i>Gigascience</i> <b>4</b> (1), 7 (2015)	23
<sup>24</sup>	11. Schlather, M., Erbe, M., Skene, F., Freudenberg, A.: miraculix: Algebraic and Statistical Functions for Genetics.	24
<sup>25</sup>	(2019). R package version 0.9.15. <a href="https://github.com/schlather/miraculix">https://github.com/schlather/miraculix</a>	25
<sup>26</sup>	12. R Core Team: <i>R: A Language and Environment for Statistical Computing</i> . R Foundation for Statistical	26
<sup>27</sup>	Computing, Vienna, Austria (2019). R Foundation for Statistical Computing. <a href="https://www.R-project.org/">https://www.R-project.org/</a>	27
<sup>28</sup>	13. Ben-Haim, G.: Practical Intel(R) AVX Optimization on 2nd generation Intel(R) Core™ Processors.	28
<sup>29</sup>	<a href="https://software.intel.com/en-us/articles/practical-intel-avx-optimization-on-2nd-generation-intel-core-processors">https://software.intel.com/en-us/articles/practical-intel-avx-optimization-on-2nd-generation-intel-core-</a>	29
<sup>30</sup>	<a href="https://software.intel.com/en-us/articles/practical-intel-avx-optimization-on-2nd-generation-intel-core-processors">processors</a> (2012)	30
<sup>31</sup>	14. Purcell, S., Neale, B., Todd-Brown, K., Thomas, L., Ferreira, M.A.R., Bender, D., Maller, J., Sklar, P., de	31
<sup>32</sup>	Bakker, P.I.W., Daly, M.J., Sham, P.C.: PLINK: a toolset for whole-genome association and population-based	32
<sup>33</sup>	linkage analysis. <i>American Journal of Human Genetics</i> <b>81</b> (3), 559–575 (2007)	33
	15. Dalke, A.: Simple benchmark harness for different popcount implementations.	
	<a href="http://dalkescientific.com/writings/diary/popcnt.cpp">http://dalkescientific.com/writings/diary/popcnt.cpp</a> (2011)	
	16. Dalke, A.: Update: Faster population counts.	
	<a href="http://www.dalkescientific.com/writings/diary/archive/2011/11/02/faster_popcount_update.html">http://www.dalkescientific.com/writings/diary/archive/2011/11/02/faster_popcount_update.html</a> (2011)	